Bust Fake News with ML & AWS

Ananth Iyer
Principal Machine Learning Engineer
LIFEOMIC

what are we doing today?

explore pre-trained model

classifier training/inference

package and deploy

what is needed

- dataset
- pretrained model
- compute for training
- hosted inference

I picked

- PolitiFact
- BERT
- p2.xlarge
- fargate

PolitiFact

https://www.politifact.com/

"Liar, Liar Pants on Fire":
A New Benchmark Dataset for Fake News Detection

William Yang Wang
Department of Computer Science
University of California, Santa Barbara
Santa Barbara, CA 93106 USA
william@cs.ucsb.edu

https://arxiv.org/abs/1705.00648

BERT

https://www.ripleys.com/weird-news/what-if-bert-was-real/

BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding

Jacob Devlin **Ming-Wei Chang** **Kenton Lee** **Kristina Toutanova**
Google AI Language
{jacobdevlin,mingweichang,kentonl,kristout}@google.com

Abstract

We introduce a new language representation model called **BERT**, which stands for **B**idirectional **E**ncoder **R**epresentations from **T**ransformers. Unlike recent language representation models (Peters et al., 2018; Radford et al., 2018), BERT is designed to pre-train deep bidirectional representations by jointly models are required to produce fine-grained output at the token-level.

There are two existing strategies for applying pre-trained language representations to downstream tasks: *feature-based* and *fine-tuning*. The feature-based approach, such as ELMo (Peters et al., 2018), uses tasks-specific architectures that include the pre-trained representations as addi-

BERT: method to model language

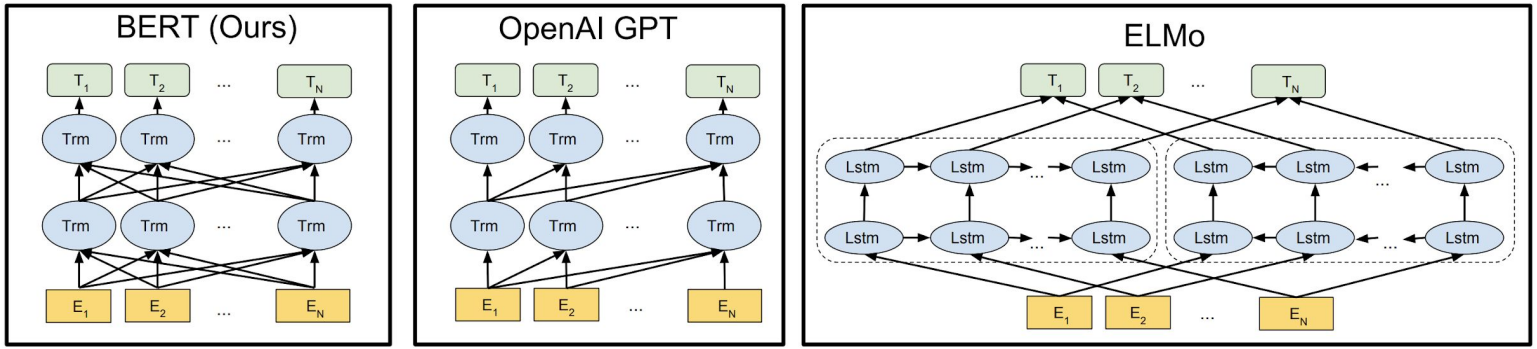

Input = `[CLS] the man went to [MASK] store [SEP]`

`he bought a gallon [MASK] milk [SEP]`

Label = `IsNext`

Input = `[CLS] the man [MASK] to the store [SEP]`

`penguin [MASK] are flight ##less birds [SEP]`

Label = `NotNext`

- 80% of the time: Replace the word with the `[MASK]` token, e.g., `my dog is hairy` → `my dog is [MASK]`
- 10% of the time: Replace the word with a random word, e.g., `my dog is hairy` → `my dog is apple`
- 10% of the time: Keep the word unchanged, e.g., `my dog is hairy` → `my dog is hairy`. The purpose of this is to bias the representation towards the actual observed word.

For the pre-training corpus we use the concatenation of BooksCorpus (800M words) (Zhu et al., 2015) and English Wikipedia (2,500M words). For Wikipedia we extract only the text passages

fargate

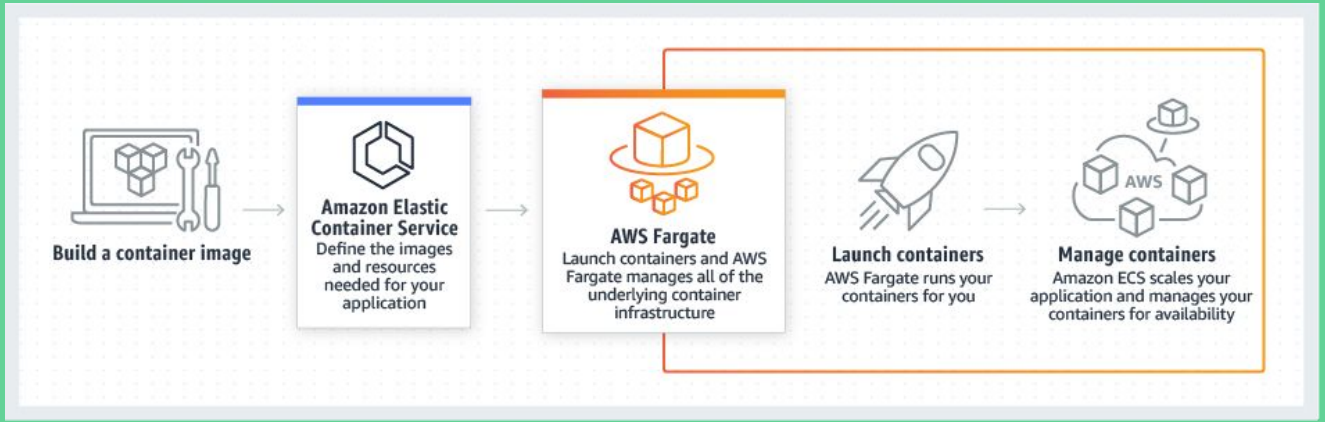

https://aws.amazon.com/fargate/

training

https://github.com/huggingface/pytorch-pretrained-BERT

```python
import torch
from torch import nn
import torch.nn.functional as F

from pytorch_pretrained_bert import BertModel
from pytorch_pretrained_bert.modeling import BertPreTrainedModel

class BertForSequenceClassificationSoftmax(BertPreTrainedModel):
    def __init__(self, config, num_labels):
        super(BertForSequenceClassificationSoftmax, self).__init__(config)
        self.num_labels = num_labels
        self.bert = BertModel(config)
        self.dropout = nn.Dropout(config.hidden_dropout_prob)
        self.classifier = nn.Linear(config.hidden_size, num_labels)
        self.apply(self.init_bert_weights)

    def forward(self, input_ids, token_type_ids=None, attention_mask=None, labels=None):
        _, pooled_output = self.bert(input_ids, token_type_ids, attention_mask)
        pooled_output = self.dropout(pooled_output)
        logits = self.classifier(pooled_output)
        probs = F.softmax(probs, dim=-1)
        if labels is not None:
            loss_fct = CrossEntropyLoss()
            loss = loss_fct(probs.view(-1, self.num_labels), labels.view(-1))
            return loss
        else:
            return probs
```

inference

```
class FakeDetector(object):
  def __init__(self, bert_dir):

    self.tokenizer = BertTokenizer.from_pretrained(bert_dir)
    self.model = BertForSequenceClassificationSoftmax.from_pretrained(bert_dir, num_labels=2)
    self.model.eval()

  def process(self, inp):
    # tokenize text and convert to indices
    tokenized_text = self.tokenizer.tokenize(inp)
    indexed_tokens = self.tokenizer.convert_tokens_to_ids(tokenized_text)

    # get sentence embedding
    probs = self.model( torch.tensor([indexed_tokens]) ).detach().numpy()[0]

    labels = ['TRUE', 'FAKE']
    prob_idx = np.argmax(probs)

    return {
      'conf': float(probs[prob_idx]),
      'label': labels[prob_idx]
    }
```

package


```python
from flask import Flask
from flask import jsonify
from flask import request
from infer.fake_detector import FakeDetector

app = Flask(__name__)

bert_dir = 'infer/model'
detector = FakeDetector(bert_dir)

@app.route('/detect', methods=['POST'])
def detect():

    content = request.get_json(silent=True)
    return jsonify({
        'result': detector.process(content['text'])
    })
```

package

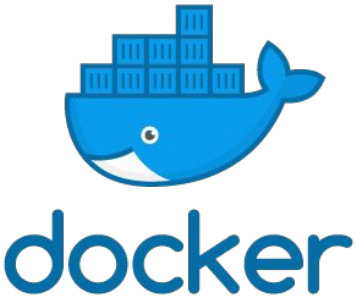


```
FROM conda/miniconda3

RUN apt-get update && apt-get install -y build-essential

RUN mkdir -p /opt/app
COPY infer /opt/app/infer
COPY bertex /opt/app/bertex
WORKDIR /opt/app

RUN conda create -n venv python=3.6 \
    && conda env update -n venv -f infer/environment.yml \
    && echo "source activate venv" >> ~/.bashrc

ENV PATH /opt/conda/envs/env/bin:$PATH
CMD /bin/bash -c "source activate venv && gunicorn -w 4 -b 0.0.0.0:8000 infer.server:app"
```

deploy

Amazon ECS

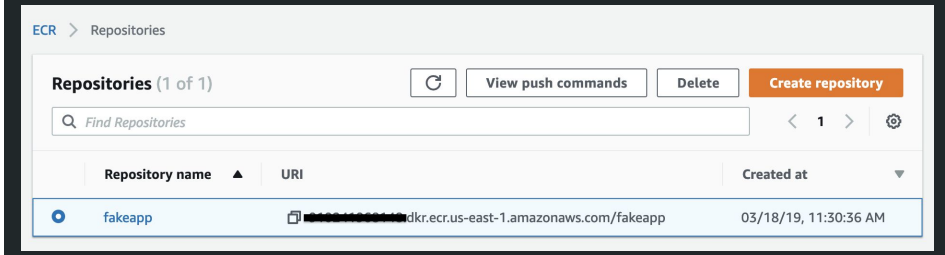

> docker build

> docker tag

> docker push

deploy

Preparing service : 10 of 10 complete

ECS resource creation	complete
Cluster fakeapp-infer-cluster	complete
Task definition fakeapp-infer-task:1	complete
Service fakeapp-infer-service	complete
Additional AWS service integrations	complete
Log group /ecs/fakeapp-infer-task	complete
CloudFormation stack EC2ContainerService-fakeapp-infer-cluster	complete
VPC vpc-077d82c77d2afc4ef	complete
Subnet 1 subnet-012080c557c3d6cea	complete
Subnet 2 subnet-084598ae661576cec	complete
Security group sg-0f7bd8a15582b3fbe	complete
Load balancer arn:aws:elasticloadbalancing:us-east-[redacted]loadbalancer/app/EC2Co-EcsEl-2X22P9OD67C0/58879591278eb330	complete

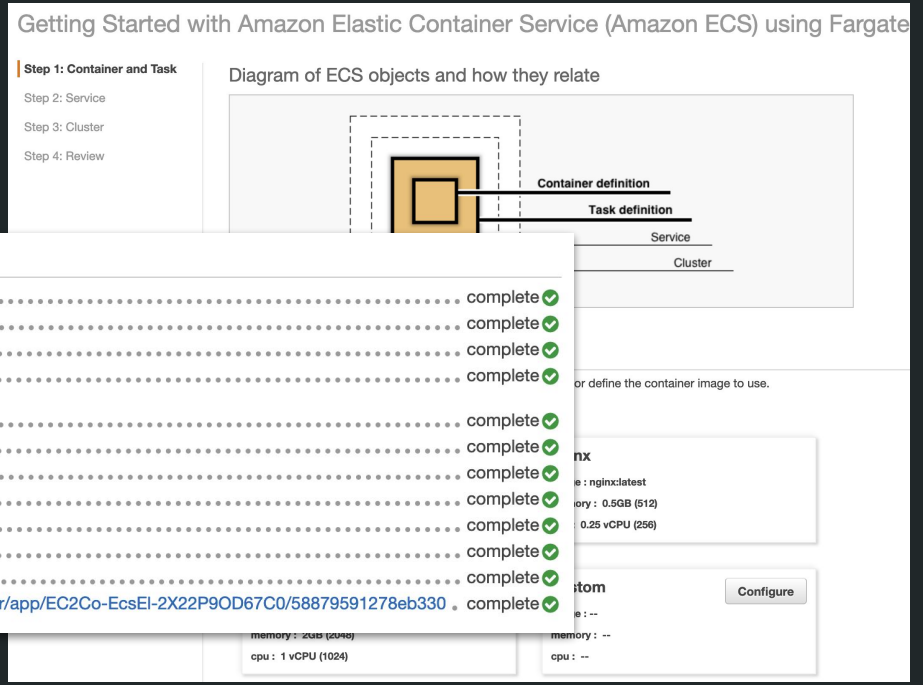

let’s give it a try

discussion

aniyer7@gmail.com

/aniyer7